

CMPE 214

GPU Architecture & Programming

# **Lecture 5.**

## **Advanced GPU Programming (2)**

Haonan Wang



SAN JOSÉ STATE  
UNIVERSITY


# Naïve Reduction Kernel

```
...  
  
__shared__ float partialSum[2*BLOCK_SIZE];  
  
unsigned int t = threadIdx.x;  
unsigned int start = 2 * blockIdx.x * blockDim.x;  
partialSum[t] = input[start + t];  
partialSum[blockDim + t] = input[start + blockDim.x + t];  
  
for (unsigned int stride = 1; stride <= blockDim.x; stride *= 2){  
    __syncthreads();  
  
    if (t % stride == 0)  
        partialSum[2 * t] += partialSum[2 * t + stride];  
}  
...
```

**How to sum up the results from different blocks?**

# A better Thread Organization

Thread 0 Thread 1 Thread 2 Thread 3



# Modified Reduction Kernel

```
...  
  
__shared__ float partialSum[2*BLOCK_SIZE];  
  
unsigned int t = threadIdx.x;  
unsigned int start = 2 * blockIdx.x * blockDim.x;  
partialSum[t] = input[start + t];  
partialSum[blockDim + t] = input[start + blockDim.x + t];  
  
for (unsigned int stride = blockDim.x/2; stride > 0; stride /= 2){  
    __syncthreads();  
  
    if (tid < stride)  
        partialSum[tid] += partialSum[tid + stride];  
}  
  
...
```

How to further optimize?

# Loop Unrolling Within A Warp

```
...  
  
for (unsigned int stride = blockDim.x/2; stride > 32; stride /= 2){  
    if (tid < stride)  
        partialSum[tid] += partialSum[tid + stride];  
    __syncthreads();  
}
```

What is good about removing the for statement?

```
if( tid < 32 ) partialSum[t] += partialSum[t + 32];  
if( tid < 16 ) partialSum[t] += partialSum[t + 16];  
if( tid < 8 ) partialSum[t] += partialSum[t + 8];  
if( tid < 4 ) partialSum[t] += partialSum[t + 4];  
if( tid < 2 ) partialSum[t] += partialSum[t + 2];  
if( tid < 1 ) partialSum[t] += partialSum[t + 1];
```

Already in lockstep within a warp

...

# Loop Unrolling in General

```
#pragma unroll
for (unsigned int stride = 32; stride > 0; stride /= 2){
    if (tid < stride) partialSum[tid] += partialSum[tid + stride];
}
```


Is the same as:

```
if( tid < 32 ) partialSum[t] += partialSum[t + 32];
if( tid < 16 ) partialSum[t] += partialSum[t + 16];
if( tid < 8 ) partialSum[t] += partialSum[t + 8];
if( tid < 4 ) partialSum[t] += partialSum[t + 4];
if( tid < 2 ) partialSum[t] += partialSum[t + 2];
if( tid < 1 ) partialSum[t] += partialSum[t + 1];
```


Unroll factor can also be specified: `#pragma unroll <UnrollFactor>`

# Warp Shuffle Instructions (1)

Threads communicate via shared memory:



Threads communicate directly?



- **T\_\_shfl\_sync(unsigned mask, T var, int srcLane, int width=warpSize);**
  - Copy from lane ID (arbitrary pattern)
- **T\_\_shfl\_up\_sync(unsigned mask, T var, unsigned int delta, int width=warpSize);**
  - Copy from delta/offset lower lane
- **T\_\_shfl\_down\_sync(unsigned mask, T var, unsigned int delta, int width=warpSize);**
  - Copy from delta/offset higher lane
- **T\_\_shfl\_xor\_sync(unsigned mask, T var, int laneMask, int width=warpSize);**
  - Copy from calculated lane ID (calculated pattern)


# Warp Shuffle Instructions (2)

- `T __shfl_sync(unsigned mask, T var, int srcLane, int width=warpSize);`
  - Copy from lane ID (arbitrary pattern)

- **Example 1:**

`Int i = threadIdx.x % 8;`


`Int j = __shfl_sync(0xffffffff, i, 2);`



- **Example 2:**

`int i = threadIdx.x % 32;`

`int j = __shfl_sync(0xffffffff, i, (threadIdx.x+1)%32, 8);`






# Warp Shuffle Instructions (3)

- **T \_\_shfl\_up\_sync(unsigned mask, T var, unsigned int delta, int width=warpSize);**
  - Copy from delta/offset lower lane

- **Example:**

```
int i = threadIdx.x % 32;
```

```
int j = __shfl_up_sync(0xffffffff, i, 2, 8);
```




- **T \_\_shfl\_down\_sync(unsigned mask, T var, unsigned int delta, int width=warpSize);**
  - Copy from delta/offset higher lane:

- **Example:**

```
int i = threadIdx.x % 32;
```

```
int j = __shfl_down_sync(0xffffffff, i, 2, 8);
```




# Warp Shuffle Instructions (4)

- `T __shfl_xor_sync(unsigned mask, T var, int laneMask, int width=warpSize);`
  - Copy from calculated lane ID

- **Example:**

`Int i = threadIdx.x % 32;`

`Int j = __shfl_xor_sync(0xffffffff, i, 1, 8);`



- An XOR (exclusive or) operation is performed between `laneMask` and the calling thread's `laneID` to determine the lane from which to copy the value.
- **Be careful with conditional code!**
  - Threads may only read data from another active thread.
  - If the target thread is inactive, the retrieved value is undefined.

# Using Warp Shuffle Instructions

```
...  
  
for (unsigned int stride = blockDim.x/2; stride > 16; stride /= 2){  
    if (tid < stride)  
        partialSum[tid] += partialSum[tid + stride];  
    __syncthreads();  
}  
  
float shuffle_sum = 0;  
if( tid < 32 ){  
    shuffle_sum = partialSum[tid];  
    for (int offset = 16; offset > 0; offset /= 2){  
        shuffle_sum += __shfl_down_sync(0xffffffff, shuffle_sum, offset);  
    }  
}  
...
```

# Spinlocks with Atomic

```
int compare_and_swap(int *value, int expected, int new_value)
{
    int original_value = *value;
    if (*value == expected)
        *value = new_value;
    return original_value;
}
```

```
while (true){
    while (compare_and_swap(&lock, 0, 1) != 0)
    {
        // busy wait
    }

    critical_region();
    lock = 0;
    noncritical_region();
}
```

# Locks in CUDA (1)

```
__global__ void myKernel(Lock lock, int *A){

    lock.lock();
    // Serial part (critical region)
    lock.unlock();
}

Int main(){
...

    Lock lock;
    myKernel<<<m, n>>> (lock, A);

...
}
```

# Locks in CUDA (2)

```
class Lock {
    Int *mutex; // lock value

public:
    Lock () {
        int state = 0;
        cudaMalloc((void**) &mutex, sizeof(int));
        cudaMemcpy(mutex, &state, sizeof(int), cudaMemcpyHostToDevice);
    }

    ~Lock () {
        cudaFree(mutex);
    }

    ...
}
```

# Locks in CUDA (3)

```
class Lock {  
...  
  
    __device__ void lock () {  
        while (atomicCAS(mutex, 0, 1) != 0);  
    }  
  
    __device__ void unlock () {  
        *mutex= 0;  
    }  
  
}
```

# Locks in CUDA (4)

```
__global__ void myKernel(Lock lock, int *A) {  
  
    lock.lock();  
    // Serial part (critical region)  
    lock.unlock();  
}
```

Any issue?



```
__global__ void myKernel(Lock lock, int *A) {  
  
    if((threadIdx.x % 32) == 0) {  
        lock.lock();  
        // Serial part (critical region)  
        lock.unlock();  
    }  
}
```



SAN JOSÉ STATE UNIVERSITY *powering* SILICON VALLEY

